

Assessment based on github pull requests	https://github.com/awroble/NLP_2025W/p	Points
Topic number / info	LLM Safety: Jailbreaks (Instruction Bypass), Prompt Injection, and Hallucination Robustness	
Team members	Alicja Charuza, Martyna Kuśmierz, Dawid Sroczyk	
Phase 2 - revised SoTA review + PoC+EDA (MAX 20 points in the course rules) - grade in orange at the end of the table; max. 1 point for each aspect in the table, if not written directly. MAX points to get actually: 20 +1 additional outcomes. EDA - 3 points, SOLUTION PLAN 3 points, Adherence to Dr Jan Sawicki's remarks - 2 points, reproducibility checklist +2 points, wrong template - "-2 points"		
CLEARLY STATED CONTRIBUTION: Table with the contribution of each team member and time assessment (workload)	no individual contributions, no workload	0
Scientific and precise language, editorial and grammar correctness	OK.	1
... meaningful references, correctly cited and reported (if exists, not from arXIV or other preprint servers)	Most are without links, or detailed information - only authors, year and title.	0
... figure and table captions easy to understand at first glance	Many figures visualising the results. As for me, they are not perfectly in line with the aesthetic standards.	1
Revised literature review (related datasets + methods)	Done, but lacks precise comparisons, too general descriptions.	0,75
Solution plan & Experimental setting - description of the experimental procedure with settings of experiments - max. 3 points	Quite ok. Models are not precisely cited.	2,5
... result analysis refers to tables and figures with results	Result analysis interesting, but lacks a bit of precision.	0,75
... adjustments of the chosen metrics (the best not only one)	The metrics are not defined, or the annotation protocol. It is rather indirectly mentioned. For the next stage, please provide multiple runs of each training and testing.	0,25
... time/memory measured?	The report shows some awareness of these issues. But it does not provide direct measurement.	0

Rebuttal or corrections for all the tips given by dr Jan Sawicki - max. 2 points	Jan's remarks are partially addressed. Jan's remarks: "- Some key metrics lack a detailed definition, e.g., "86.1% success rate" seems undefined. - In plus: clear structure and well-defined areas of research/project"	1
Presentation (26.11.2025, or next Wednesday) - mandatory, not graded here	delivered, not provided the presentation file.	
EDA - comparison to different datasets (depends on the project topic) - max. 3 points	Done, however, with no direct comparisons of differences and lacks precision, e.g., categories are listed, and only the number of observations is provided.	1,5
Preliminary results (machine learning): results with analysis, comparison to different settings or models (depends on the project topic)	Not very precise, but interesting.	1
Reasonably clean code is delivered - clean, reproducible code	Mostly clean, only the last category - prompt injection - remains with only two notebooks.	0,75
Readme to understand the code, code structure (folders)	The same comment as above, only one category has readme file.	0,5
Provided processed datasets!	The same comment as above. If I have contribution table I would correct individually my assessment.	0,5
Additional outcomes - pre-processed datasets, model parameters / wrong template for the report (-1 point); reproducibility checklist at this stage counts for 2 points.	It is very unclear what the reason is for one part to be so underdeveloped. Please resolve this issue for the next stage of the final project. Also, I suspect an unfair work contribution. I want to assess the project fairly, but I do not have contribution table. Please, provide this, and amend the part of source code.	0
When received	Dec 17 2025	
Delayed ?	NO	
SUM (Score)	11,50	